

Eye Disease Classifier

April 3, 2025

```
[1]: # Import necessary libraries
import os
import numpy as np
import pandas as pd
import tensorflow as tf
import keras as keras
import seaborn as sns
import matplotlib.pyplot as plt

# Import other libraries
from keras import layers, models
from transformers import TFFViTModel
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import image_dataset_from_directory
```

```
[2]: # Physical device configuration
physical_devices = tf.config.list_physical_devices('GPU')
for gpu in physical_devices:
    tf.config.experimental.set_memory_growth(gpu, True)

print("Num GPUs Available: ", len(physical_devices)) # 1 or 0 (1 if GPU is available)
```

Num GPUs Available: 1

```
[3]: # Constants Variables
EPOCHS = 100

DATA_SET = 'dataset'
IMAGE_SIZE = 224
BATCH_SIZE = 16
CHANNELS = 3
SEED = 123
```

```
[4]: # Load the dataset
dataset = image_dataset_from_directory(
    DATA_SET,
    image_size=(IMAGE_SIZE, IMAGE_SIZE),
```

```

    batch_size=BATCH_SIZE,
    shuffle=True,
    seed=SEED,
)

```

Found 16242 files belonging to 10 classes.

```

[5]: # Print the class names
class_names = dataset.class_names

# Get the class names
print("Dataset Classes:")
for i, class_name in enumerate(class_names):
    print(f"{i + 1}. {class_name.replace('_', ' ')}")

```

Dataset Classes:

1. Central Serous Chorioretinopathy
2. Diabetic Retinopathy
3. Disc Edema
4. Glaucoma
5. Healthy
6. Macular Scar
7. Myopia
8. Pterygium
9. Retinal Detachment
10. Retinitis Pigmentosa

```

[6]: # Calculate dataset size
dataset_size = tf.data.experimental.cardinality(dataset).numpy()
total_samples = dataset_size * 16 # assuming batch_size = 16

print(f"Total samples: {total_samples}")

```

Total samples: 16256

```

[7]: # Length of the dataset
dataset_length = len(dataset)
print(f"Dataset Length: {dataset_length}")


```

Dataset Length: 1016

```

[8]: # Visualize the dataset
plt.figure(figsize=(10, 10))
for iimage_batch, labels_batch in dataset.take(1):
    for i in range(16):
        ax = plt.subplot(4, 4, i + 1)
        plt.imshow(iimage_batch[i].numpy().astype("uint8"))
        plt.title(class_names[labels_batch[i]])
        plt.axis("off")

```



```
[9]: # Split ratios
train_ratio, val_ratio = 0.7, 0.15

# Calculate sizes
train_size = int(train_ratio * dataset_size)
val_size = int(val_ratio * dataset_size)
test_size = dataset_size - train_size - val_size

# Shuffle with full buffer size
dataset = dataset.shuffle(buffer_size=dataset_size, seed=123)
```

```
[10]: # Split the dataset
train_dataset = dataset.take(train_size)
val_test_dataset = dataset.skip(train_size)

val_dataset = val_test_dataset.take(val_size)
test_dataset = val_test_dataset.skip(val_size)
```

```
[11]: # Add performance optimization
AUTOTUNE = tf.data.AUTOTUNE
train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
val_dataset = val_dataset.prefetch(buffer_size=AUTOTUNE)
test_dataset = test_dataset.prefetch(buffer_size=AUTOTUNE)
```

```
[12]: # Print dataset sizes
print("Train batches:", train_size)
print("Validation batches:", val_size)
print("Test batches:", test_size)
```

```
Train batches: 711
Validation batches: 152
Test batches: 153
```

```
[ ]: # Load the pre-trained ViT model
base_model = TFViTModel.from_pretrained("google/vit-large-patch16-224")

# Define the Class for the model
classes = 10

# Define the ViT layer
def vit_layer(x):
    return base_model({"pixel_values": x}).last_hidden_state[:, 0, :]

# Define the model
resize_rescale = keras.Sequential([
    keras.layers.Resizing(224, 224),
    keras.layers.Rescaling(1. / 255),
    keras.layers.Permute((3, 1, 2))
])

inputs = keras.layers.Input(shape=(224, 224, 3))
x = resize_rescale(inputs)
x = keras.layers.Lambda(vit_layer)(x)
outputs = keras.layers.Dense(classes, activation="softmax")(x)

model = keras.Model(inputs, outputs)
```

```
[14]: # Summary (optional)
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
sequential (Sequential)	(None, 3, 224, 224)	0
lambda (Lambda)	(None, 1024)	0
dense (Dense)	(None, 10)	10250

Total params: 10,250
 Trainable params: 10,250
 Non-trainable params: 0

```
[15]: # Early stopping and learning rate reduction callbacks
early_stop = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

# Learning rate reduction callback
reduce_lr = tf.keras.callbacks.ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.5,
    patience=3,
    verbose=1,
    min_lr=1e-6
)
```

```
[16]: base_model.trainable = False # Freeze ViT base model for 1st phase of training
```

```
[17]: # compile the model
model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

# then fit
history = model.fit(
    train_dataset,
```

```
validation_data=val_dataset,  
epochs=EPOCHS,  
callbacks=[early_stop, reduce_lr]  
)
```

Epoch 1/100

711/711 [=====] - 215s 251ms/step - loss: 1.0500 -
accuracy: 0.6296 - val_loss: 0.8089 - val_accuracy: 0.7056

Epoch 2/100

711/711 [=====] - 192s 245ms/step - loss: 0.7449 -
accuracy: 0.7293 - val_loss: 0.6576 - val_accuracy: 0.7558

Epoch 3/100

711/711 [=====] - 192s 244ms/step - loss: 0.6794 -
accuracy: 0.7493 - val_loss: 0.5913 - val_accuracy: 0.7961

Epoch 4/100

711/711 [=====] - 191s 243ms/step - loss: 0.6261 -
accuracy: 0.7706 - val_loss: 0.5881 - val_accuracy: 0.7907

Epoch 5/100

711/711 [=====] - 191s 243ms/step - loss: 0.5914 -
accuracy: 0.7802 - val_loss: 0.5394 - val_accuracy: 0.7915

Epoch 6/100

711/711 [=====] - 191s 243ms/step - loss: 0.5731 -
accuracy: 0.7833 - val_loss: 0.5405 - val_accuracy: 0.8014

Epoch 7/100

711/711 [=====] - 191s 243ms/step - loss: 0.5522 -
accuracy: 0.7941 - val_loss: 0.5005 - val_accuracy: 0.8122

Epoch 8/100

711/711 [=====] - 191s 244ms/step - loss: 0.5365 -
accuracy: 0.7987 - val_loss: 0.5122 - val_accuracy: 0.8121

Epoch 9/100

711/711 [=====] - 191s 244ms/step - loss: 0.5192 -
accuracy: 0.8034 - val_loss: 0.5010 - val_accuracy: 0.7998

Epoch 10/100

711/711 [=====] - 191s 243ms/step - loss: 0.5086 -
accuracy: 0.8102 - val_loss: 0.4910 - val_accuracy: 0.8220

Epoch 11/100

711/711 [=====] - 191s 243ms/step - loss: 0.4905 -
accuracy: 0.8117 - val_loss: 0.4708 - val_accuracy: 0.8220

Epoch 12/100

711/711 [=====] - 191s 243ms/step - loss: 0.4877 -
accuracy: 0.8143 - val_loss: 0.4312 - val_accuracy: 0.8420

Epoch 13/100

711/711 [=====] - 191s 244ms/step - loss: 0.4685 -
accuracy: 0.8220 - val_loss: 0.4682 - val_accuracy: 0.8178

Epoch 14/100

711/711 [=====] - 191s 244ms/step - loss: 0.4648 -
accuracy: 0.8228 - val_loss: 0.4264 - val_accuracy: 0.8376

Epoch 15/100

711/711 [=====] - 191s 244ms/step - loss: 0.4611 - accuracy: 0.8267 - val_loss: 0.4274 - val_accuracy: 0.8450
Epoch 16/100
711/711 [=====] - 191s 244ms/step - loss: 0.4625 - accuracy: 0.8226 - val_loss: 0.4333 - val_accuracy: 0.8376
Epoch 17/100
711/711 [=====] - 191s 244ms/step - loss: 0.4508 - accuracy: 0.8292 - val_loss: 0.4539 - val_accuracy: 0.8187

Epoch 00017: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.
Epoch 18/100
711/711 [=====] - 191s 243ms/step - loss: 0.4152 - accuracy: 0.8380 - val_loss: 0.3815 - val_accuracy: 0.8516
Epoch 19/100
711/711 [=====] - 191s 243ms/step - loss: 0.4127 - accuracy: 0.8456 - val_loss: 0.4065 - val_accuracy: 0.8404
Epoch 20/100
711/711 [=====] - 191s 243ms/step - loss: 0.4074 - accuracy: 0.8435 - val_loss: 0.3911 - val_accuracy: 0.8540
Epoch 21/100
711/711 [=====] - 191s 244ms/step - loss: 0.4013 - accuracy: 0.8485 - val_loss: 0.4163 - val_accuracy: 0.8421

Epoch 00021: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.
Epoch 22/100
711/711 [=====] - 191s 243ms/step - loss: 0.3927 - accuracy: 0.8508 - val_loss: 0.3803 - val_accuracy: 0.8540
Epoch 23/100
711/711 [=====] - 191s 244ms/step - loss: 0.3875 - accuracy: 0.8534 - val_loss: 0.3627 - val_accuracy: 0.8676
Epoch 24/100
711/711 [=====] - 191s 244ms/step - loss: 0.3817 - accuracy: 0.8566 - val_loss: 0.3657 - val_accuracy: 0.8709
Epoch 25/100
711/711 [=====] - 191s 243ms/step - loss: 0.3855 - accuracy: 0.8535 - val_loss: 0.3781 - val_accuracy: 0.8479
Epoch 26/100
711/711 [=====] - 191s 243ms/step - loss: 0.3838 - accuracy: 0.8519 - val_loss: 0.3971 - val_accuracy: 0.8405

Epoch 00026: ReduceLROnPlateau reducing learning rate to 0.0001250000059371814.
Epoch 27/100
711/711 [=====] - 191s 243ms/step - loss: 0.3730 - accuracy: 0.8605 - val_loss: 0.3648 - val_accuracy: 0.8660
Epoch 28/100
711/711 [=====] - 191s 244ms/step - loss: 0.3704 - accuracy: 0.8624 - val_loss: 0.3753 - val_accuracy: 0.8573
Epoch 29/100

711/711 [=====] - 191s 244ms/step - loss: 0.3763 -
accuracy: 0.8572 - val_loss: 0.3511 - val_accuracy: 0.8729

Epoch 30/100

711/711 [=====] - 191s 243ms/step - loss: 0.3745 -
accuracy: 0.8576 - val_loss: 0.3759 - val_accuracy: 0.8623

Epoch 31/100

711/711 [=====] - 191s 244ms/step - loss: 0.3712 -
accuracy: 0.8603 - val_loss: 0.3614 - val_accuracy: 0.8602

Epoch 32/100

711/711 [=====] - 191s 243ms/step - loss: 0.3668 -
accuracy: 0.8611 - val_loss: 0.3615 - val_accuracy: 0.8668

Epoch 00032: ReduceLROnPlateau reducing learning rate to 6.25000029685907e-05.

Epoch 33/100

711/711 [=====] - 191s 244ms/step - loss: 0.3673 -
accuracy: 0.8615 - val_loss: 0.3713 - val_accuracy: 0.8618

Epoch 34/100

711/711 [=====] - 191s 243ms/step - loss: 0.3595 -
accuracy: 0.8660 - val_loss: 0.3590 - val_accuracy: 0.8631

Epoch 35/100

711/711 [=====] - 191s 244ms/step - loss: 0.3663 -
accuracy: 0.8624 - val_loss: 0.3896 - val_accuracy: 0.8544

Epoch 00035: ReduceLROnPlateau reducing learning rate to 3.125000148429535e-05.

Epoch 36/100

711/711 [=====] - 191s 243ms/step - loss: 0.3600 -
accuracy: 0.8627 - val_loss: 0.3658 - val_accuracy: 0.8680

Epoch 37/100

711/711 [=====] - 191s 244ms/step - loss: 0.3649 -
accuracy: 0.8630 - val_loss: 0.3511 - val_accuracy: 0.8697

Epoch 38/100

711/711 [=====] - 191s 244ms/step - loss: 0.3625 -
accuracy: 0.8651 - val_loss: 0.3479 - val_accuracy: 0.8672

Epoch 39/100

711/711 [=====] - 191s 244ms/step - loss: 0.3628 -
accuracy: 0.8631 - val_loss: 0.3594 - val_accuracy: 0.8664

Epoch 40/100

711/711 [=====] - 191s 244ms/step - loss: 0.3634 -
accuracy: 0.8632 - val_loss: 0.3743 - val_accuracy: 0.8553

Epoch 41/100

711/711 [=====] - 191s 244ms/step - loss: 0.3621 -
accuracy: 0.8632 - val_loss: 0.3685 - val_accuracy: 0.8643

Epoch 00041: ReduceLROnPlateau reducing learning rate to 1.5625000742147677e-05.

Epoch 42/100

711/711 [=====] - 191s 244ms/step - loss: 0.3618 -
accuracy: 0.8649 - val_loss: 0.3794 - val_accuracy: 0.8573

Epoch 43/100


```

711/711 [=====] - 191s 244ms/step - loss: 0.3612 -
accuracy: 0.8621 - val_loss: 0.3609 - val_accuracy: 0.8697
Epoch 44/100
711/711 [=====] - 191s 244ms/step - loss: 0.3683 -
accuracy: 0.8620 - val_loss: 0.3442 - val_accuracy: 0.8713
Epoch 45/100
711/711 [=====] - 191s 244ms/step - loss: 0.3583 -
accuracy: 0.8676 - val_loss: 0.3579 - val_accuracy: 0.8664
Epoch 46/100
711/711 [=====] - 191s 244ms/step - loss: 0.3608 -
accuracy: 0.8653 - val_loss: 0.3486 - val_accuracy: 0.8713
Epoch 47/100
711/711 [=====] - 191s 244ms/step - loss: 0.3585 -
accuracy: 0.8662 - val_loss: 0.3665 - val_accuracy: 0.8618

Epoch 00047: ReduceLROnPlateau reducing learning rate to 7.812500371073838e-06.
Epoch 48/100
711/711 [=====] - 192s 244ms/step - loss: 0.3567 -
accuracy: 0.8672 - val_loss: 0.3602 - val_accuracy: 0.8664
Epoch 49/100
711/711 [=====] - 191s 244ms/step - loss: 0.3645 -
accuracy: 0.8640 - val_loss: 0.3587 - val_accuracy: 0.8643
Epoch 50/100
711/711 [=====] - 191s 244ms/step - loss: 0.3574 -
accuracy: 0.8660 - val_loss: 0.3608 - val_accuracy: 0.8655

Epoch 00050: ReduceLROnPlateau reducing learning rate to 3.906250185536919e-06.
Epoch 51/100
711/711 [=====] - 191s 244ms/step - loss: 0.3611 -
accuracy: 0.8642 - val_loss: 0.3534 - val_accuracy: 0.8681
Epoch 52/100
711/711 [=====] - 191s 244ms/step - loss: 0.3589 -
accuracy: 0.8673 - val_loss: 0.3574 - val_accuracy: 0.8623
Epoch 53/100
711/711 [=====] - 191s 244ms/step - loss: 0.3578 -
accuracy: 0.8650 - val_loss: 0.3782 - val_accuracy: 0.8606

Epoch 00053: ReduceLROnPlateau reducing learning rate to 1.9531250927684596e-06.
Epoch 54/100
711/711 [=====] - 191s 244ms/step - loss: 0.3590 -
accuracy: 0.8651 - val_loss: 0.3645 - val_accuracy: 0.8606

```

```

[18]: # Save the model
model.save("model/eyes_diseases.h5")
model.save("model/eyes_diseases.keras")


```

```
[19]: # Create figure and axis objects with a single subplot
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 10))

# Plot training & validation accuracy values
ax1.plot(history.history['accuracy'], label='Training')
ax1.plot(history.history['val_accuracy'], label='Validation')
ax1.set_title('Model Accuracy')
ax1.set_xlabel('Epoch')
ax1.set_ylabel('Accuracy')
ax1.legend(loc='lower right')
ax1.grid(True)

# Plot training & validation loss values
ax2.plot(history.history['loss'], label='Training')
ax2.plot(history.history['val_loss'], label='Validation')
ax2.set_title('Model Loss')
ax2.set_xlabel('Epoch')
ax2.set_ylabel('Loss')
ax2.legend(loc='upper right')
ax2.grid(True)

# Adjust the layout and display the plot
plt.tight_layout()
plt.show()
```



```
[24]: from sklearn.metrics import classification_report

def evaluate_and_report_model(model, test_dataset, class_names):
```

```

loss, acc = model.evaluate(test_dataset)
print(f"Test Loss: {loss:.4f}")
print(f"Test Accuracy: {acc:.4f}")

y_true = []
y_pred = []

for images, labels in test_dataset:
    preds = model.predict(images)
    y_true.extend(labels.numpy())
    y_pred.extend(np.argmax(preds, axis=1))

y_true = np.array(y_true)
y_pred = np.array(y_pred)

print("\n=== Classification Report ===")
print(classification_report(y_true, y_pred, target_names=class_names))

evaluate_and_report_model(model, test_dataset, class_names)

```

153/153 [=====] - 46s 185ms/step - loss: 0.3512 - accuracy: 0.8668

Test Loss: 0.3512

Test Accuracy: 0.8668

=== Classification Report ===

	precision	recall	f1-score	support
Central Serous Chorioretinopathy	0.83	0.81	0.82	111
Diabetic Retinopathy	0.96	0.96	0.96	536
Disc Edema	0.96	0.97	0.96	119
Glaucoma	0.75	0.78	0.76	389
Healthy	0.79	0.81	0.80	391
Macular Scar	0.83	0.77	0.80	294
Myopia	0.87	0.86	0.87	359
Pterygium	1.00	1.00	1.00	18
Retinal Detachment	1.00	0.98	0.99	97
Retinitis Pigmentosa	0.96	0.97	0.97	134
accuracy			0.86	2448
macro avg	0.89	0.89	0.89	2448
weighted avg	0.87	0.86	0.86	2448

```

[29]: # Predictions using the model
for image_batch, labels_batch in test_dataset.take(1): # Take one batch from
    ↪ test dataset

```

```

ps = model.predict(image_batch) # Get model predictions for the batch
images = image_batch.numpy().astype("uint8") # Convert images to numpy
↪array
labels = labels_batch.numpy() # Convert labels to numpy array

plt.figure(figsize=(15, 15)) # Create a figure with specified size
for i in range(12): # Loop through first 12 images
    ax = plt.subplot(4, 4, i + 1) # Create a 4x4 subplot grid
    plt.imshow(images[i]) # Display the image

    # Get prediction confidence
    confidence = ps[i][ps[i].argmax()] * 100

    # Check if prediction is correct
    is_correct = labels[i] == ps[i].argmax()
    color = "green" if is_correct else "red"

    # Create title with actual label, predicted label, and confidence
    title = f"Actual: {class_names[labels[i]]}\nPredicted:
↪{class_names[ps[i].argmax()]}"
    plt.title(title, color=color) # Show title with color
    plt.axis("off") # Hide axes

```

